

# Agenda – Day 1

Time	Title	Presenter
9:00-9:10	Welcome	Frank McKenna
9:10-10:30	Debugging	Wael Elhaddad
10:30-11:00	PI / Exercises #3!	
11:00-11:15	Parallel Machines	Frank McKenna
11:15-12:00	Parallel Programming With MPI	Frank McKenna
12:00-1:00	LUNCH	
1:00-2:30	EXERCISE	YOU
2:30-3:00	Parallel Programming with OpenMP	Frank McKenna
3:00-4:30	EXERCISE	YOU
4:30-5:00	Load Balancing	Frank McKenna
Day 3	Abstraction, More C & Introducing C++	
Day 4	User Interface Design & Qt	
Day 5	SimCenter & Cloud Computing	



# *Parallel Programming*

## *Frank McKenna*



Berkeley  
UNIVERSITY OF CALIFORNIA

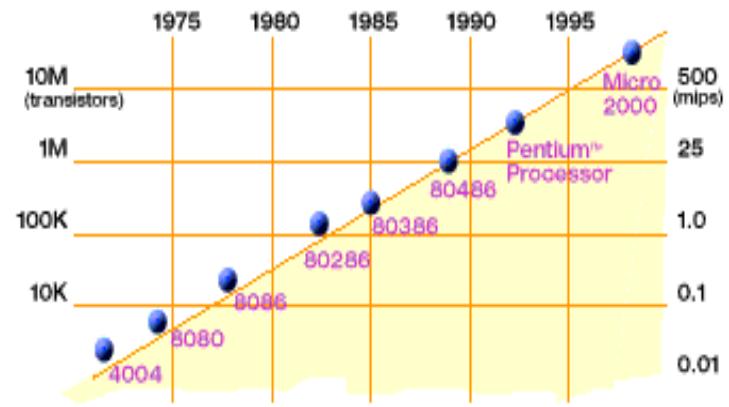
# Outline

- Parallel Machines & Parallel Machine Models
- Parallel Programming
  - Message Passing With MPI
  - Shared Memory Programming with OpenMP

**Ignoring co-processors and GPUs**

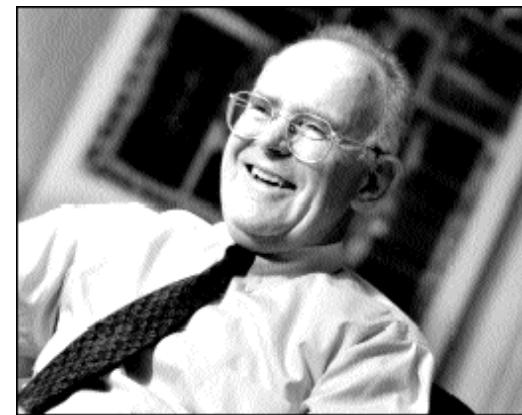
many slides source: CS267, Jim Demmel

# Why is Parallel Programming Important



2X transistors/Chip Every 1.5 years  
Called "Moore's Law"

**Microprocessors have become smaller, denser, and more powerful.**

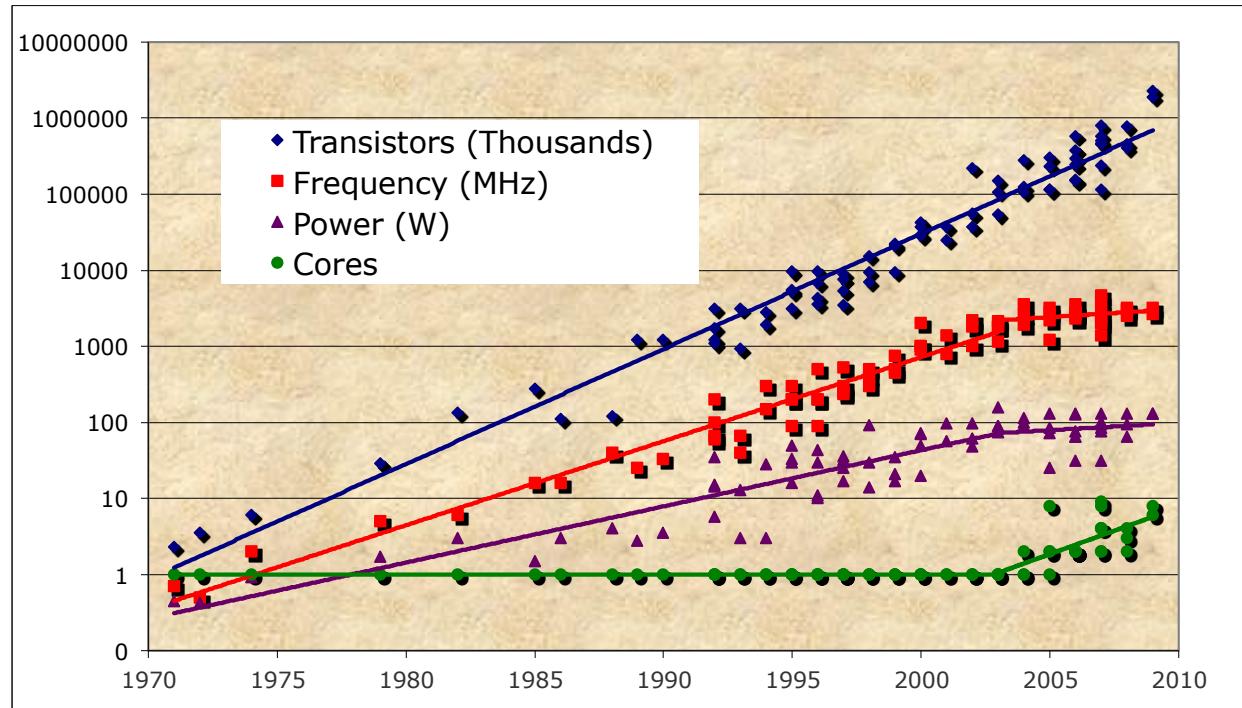


**Gordon Moore (co-founder of Intel) predicted in 1965 that the transistor density of semiconductor chips would double roughly every 18 months.**

Slide source: Jack Dongarra

source: CS267, Jim Demmel

# Revolution in Processors



- Chip density is continuing increase  $\sim 2x$  every 2 years
- Clock speed is not
- Number of processor cores may double instead
- Power is under control, no longer growing

# Multiple Cores!

What does it mean for Programmers  
“The Free Lunch is Over” Herb Sutter

- Up until 2003 programmers had been relying on Hardware to make their programs go faster. No longer. They had to start programming again!
- Performance now comes from Software
- To be fast and utilize the resources, Software must run in parallel, that is it must run on multiple cores at same time.

# How many cores?



1,736 Intel Xeon Skylake nodes, each with 48 cores + 192GB of RAM  
4,200 Intel Knights Landing nodes, each with 68 cores + 96GB of DDR RAM



1 Intel i7 node, 6 cores + 16GB RAM



Apple A11, 6 cores + 64GB RAM

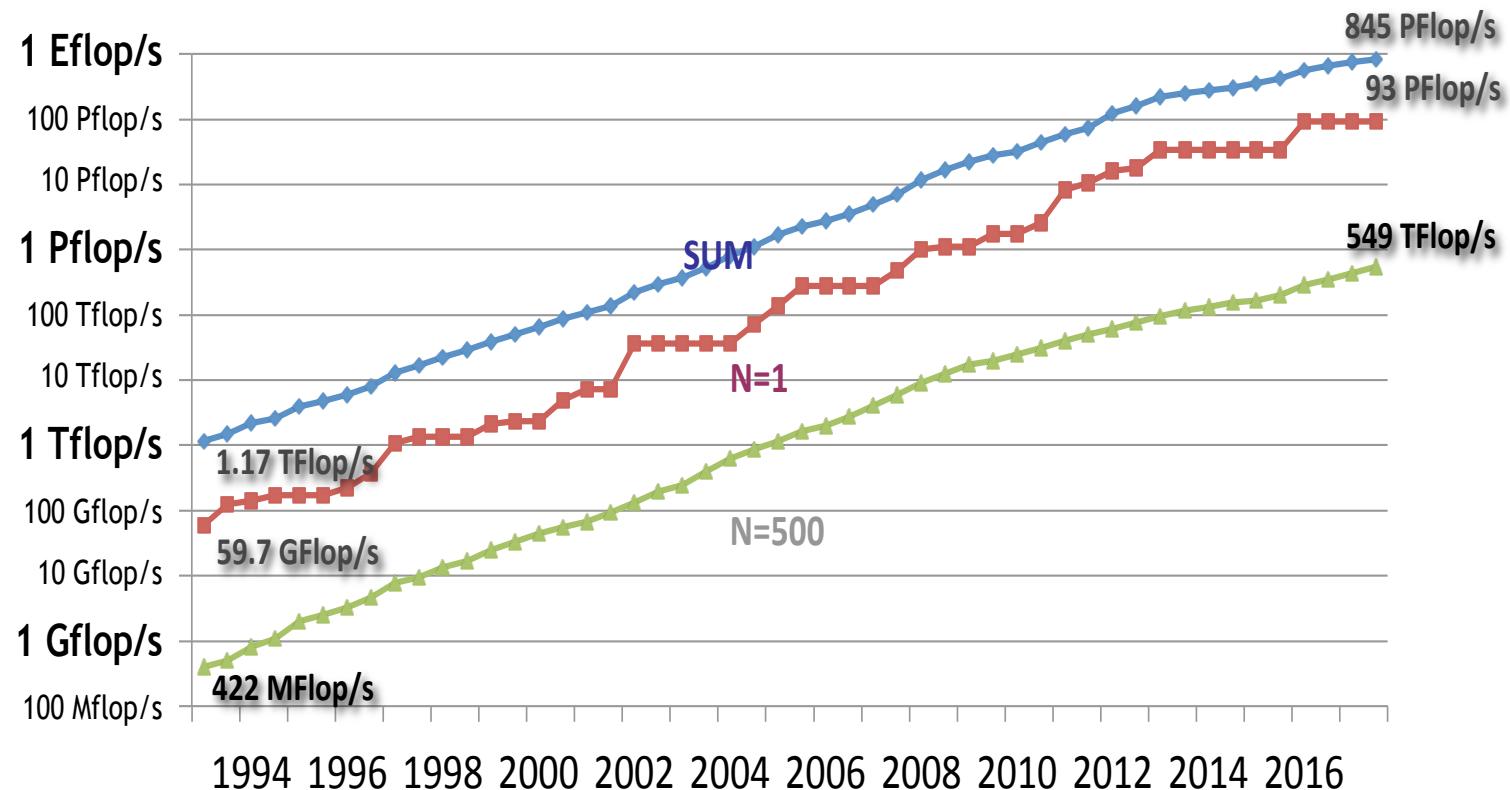
# How Big & Fast!



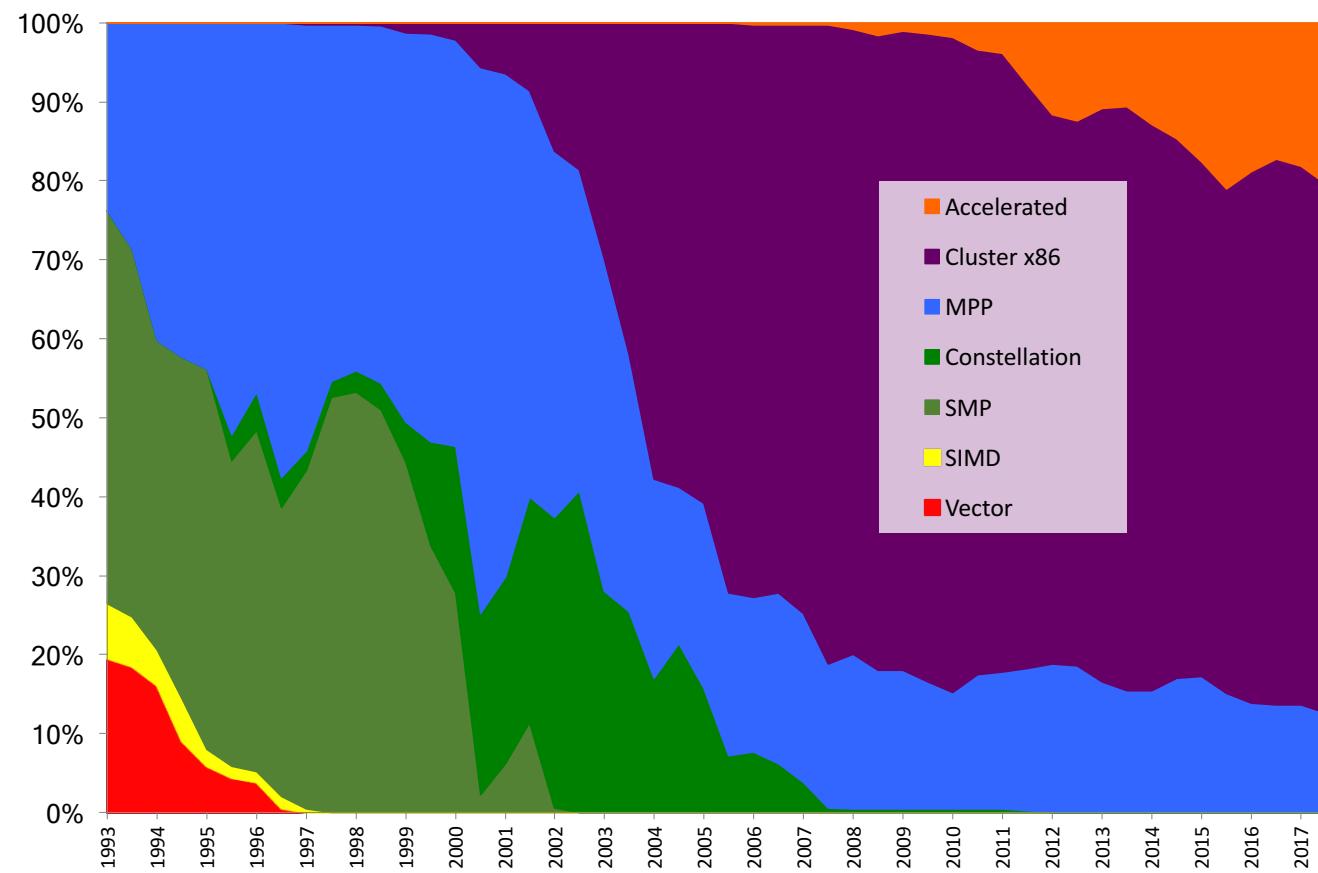
Rmax of Linpack  
Solve Ax = b

Rank	Site	System	Cores	Rmax (TFlop/s)	Rpeak (TFlop/s)	Power (kW)
1	DOE/SC/Oak Ridge National Laboratory United States	<b>Summit</b> - IBM Power System AC922, IBM POWER9 22C 3.07GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR Infiniband IBM	2,282,544	122,300.0	187,659.3	8,806
2	National Supercomputing Center in Wuxi China	<b>Sunway TaihuLight</b> - Sunway MPP, Sunway SW26010 260C 1.45GHz, Sunway NRCPC	10,649,600	93,014.6	125,435.9	15,371
3	DOE/NNSA/LLNL United States	<b>Sierra</b> - IBM Power System S922LC, IBM POWER9 22C 3.1GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR Infiniband IBM	1,572,480	71,610.0	119,193.6	
4	National Super Computer Center in Guangzhou China	<b>Tianhe-2A</b> - TH-IVB-FEP Cluster, Intel Xeon E5-2692v2 12C 2.2GHz, TH Express-2, Matrix-2000 NUDT	4,981,760	61,444.5	100,678.7	18,482
5	National Institute of Advanced Industrial Science and Technology (AIST) Japan	<b>AI Bridging Cloud Infrastructure (ABCi)</b> - PRIMERGY CX2550 M4, Xeon Gold 6148 20C 2.4GHz, NVIDIA Tesla V100 SXM2, Infiniband EDR Fujitsu	391,680	19,880.0	32,576.6	1,649
15	Texas Advanced Computing Center/Univ. of Texas United States	<b>Stampede2</b> - PowerEdge C6320P/C6420, Intel Xeon Phi 7250 68C 1.4GHz/Platinum 8160, Intel Omni-Path Dell EMC	367,024	10,680.7	18,309.2	

# Performance Development (2018)



# From Vector Supercomputers to Massively Parallel Accelerator Systems



# Moore's Law reinterpreted

- Number of cores per chip can double every two years
- Clock speed will not increase (possibly decrease)
- Need to deal with systems with millions of concurrent threads
- Need to deal with inter-chip parallelism as well as intra-chip parallelism

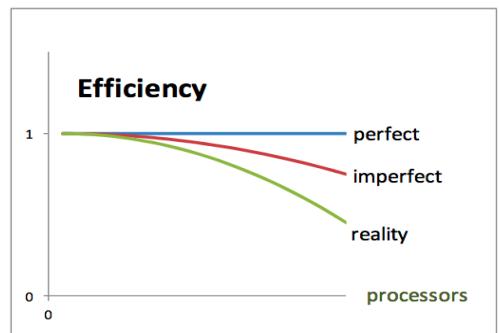
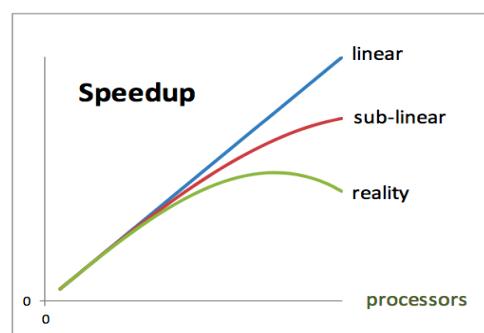
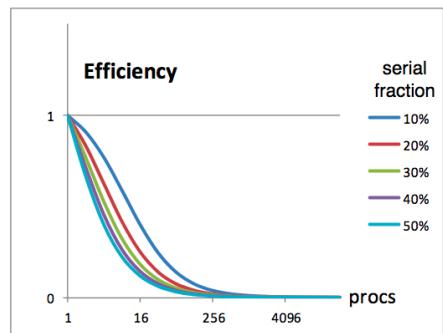
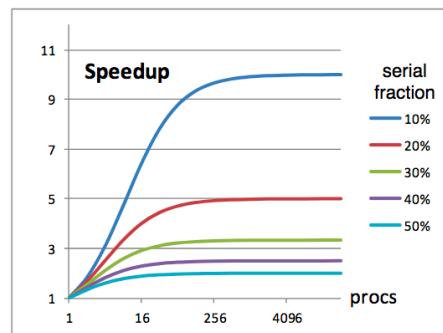
# Can All Programs Be Made to Run Faster?

- Suppose only part of an application can run in parallel
- Amdahl's law
  - let  $s$  be the fraction of work done sequentially, so  $(1-s)$  is fraction parallelizable
  - $P$  = number of processors

$$\begin{aligned}\text{Speedup}(P) &= \text{Time}(1)/\text{Time}(P) \\ &\leq 1/(s + (1-s)/P) \\ &\leq 1/s\end{aligned}$$

**QUIZ: if 10% of program  
is sequential, What is the  
maximum speedup I Can  
obtain?**

- Even if the parallel part speeds up perfectly performance is limited by the sequential part
- Top500 list: currently fastest machine has  $P \sim 2.2M$ ; Stampede2 has 367,000



Source: Doug James, TACC

# This Does not Take into Account Overhead of Parallelism

- Parallelism overheads include:
  - cost of starting a thread or process
  - cost of communicating shared data
  - cost of synchronizing
  - extra (redundant) computation
- Each of these can be in the range of milliseconds (=millions of flops) on some systems
- Tradeoff: Algorithm needs sufficiently large units of work to run fast in parallel (i.e. large granularity), but not so large that there is not enough parallel work

source: CS267, Jim Demmel

# Load Imbalance

- Load imbalance is the time that some processors in the system are idle due to
  - insufficient parallelism (during that phase)
  - unequal size tasks
- Examples of the latter
  - adapting to “interesting parts of a domain”
  - tree-structured computations
  - fundamentally unstructured problems
- Algorithm needs to balance load
  - Sometimes can determine work load, divide up evenly, before starting
    - “Static Load Balancing”
  - Sometimes work load changes dynamically, need to rebalance dynamically
    - “Dynamic Load Balancing,” eg work-stealing

# Improving Real Performance

**Peak Performance grows exponentially,  
a la Moore's Law**

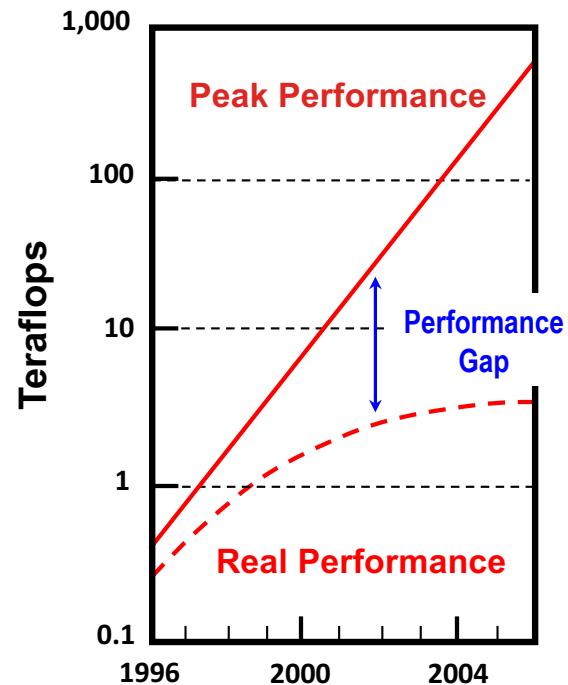
- In 1990's, peak performance increased 100x; in 2000's, it will increase 1000x

**But efficiency (the performance relative to the hardware peak) has declined**

- was 40-50% on the vector supercomputers of 1990s
- now as little as 5-10% on parallel supercomputers of today

**Close the gap through ...**

- Mathematical methods and algorithms that achieve high performance on a single processor and scale to thousands of processors
- More efficient programming models and tools for massively parallel supercomputers

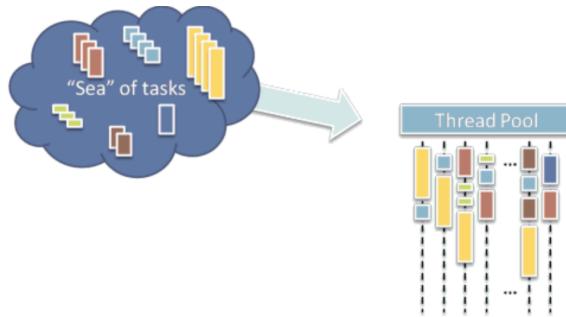


# Art of Programming - II

- To take a problem, and continually break it down into a series of smaller ideally concurrent tasks until ultimately these tasks become a series of small specific individual instructions.
- Mindful of the architecture on which the program will run, identify those tasks which can be run concurrently and map those tasks onto the processing units of the target architecture.

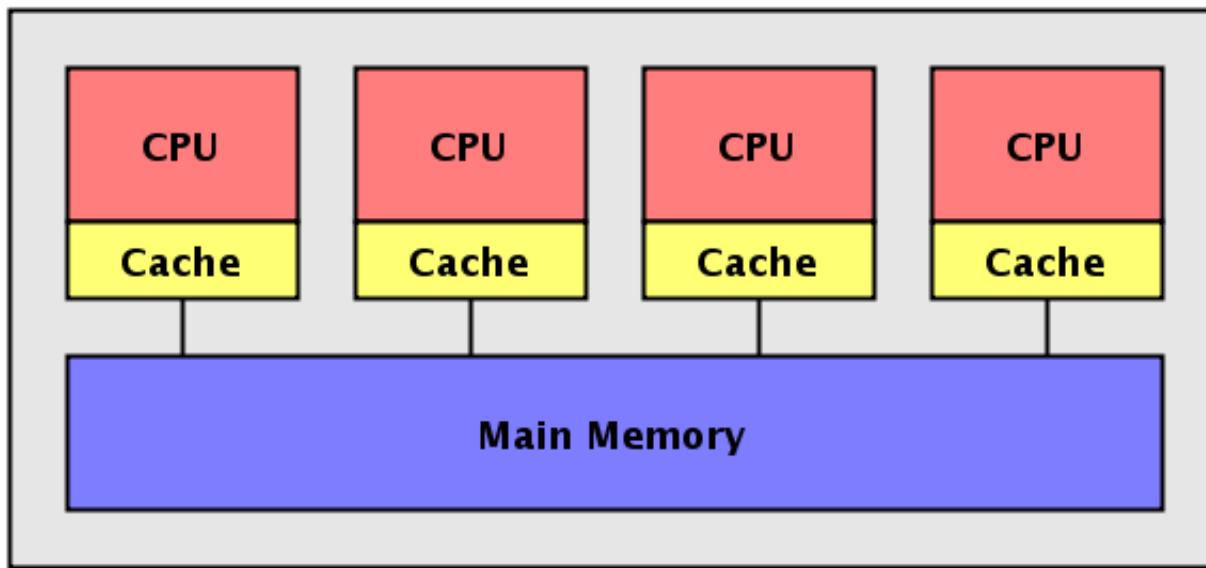
# Considerations for Parallel Programming:

- Finding enough parallelism (Amdahl's Law)
- Granularity – how big should each parallel task be
- Locality – moving data costs more than arithmetic
- Load balance – don't want 1K processors to wait for one slow one
- Coordination and synchronization – sharing data safely
- Performance modeling/debugging/tuning
- Where to put the task,



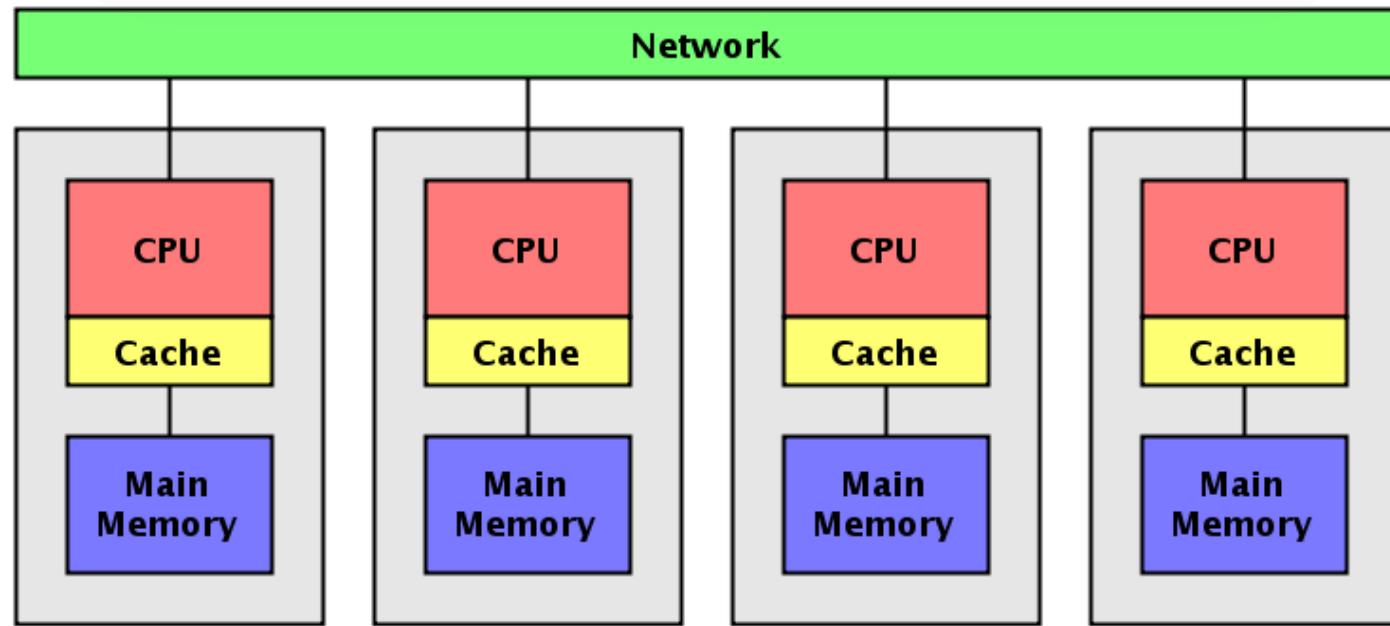
All of these things makes parallel programming even harder than sequential programming.

# Simplified Parallel Machine Models



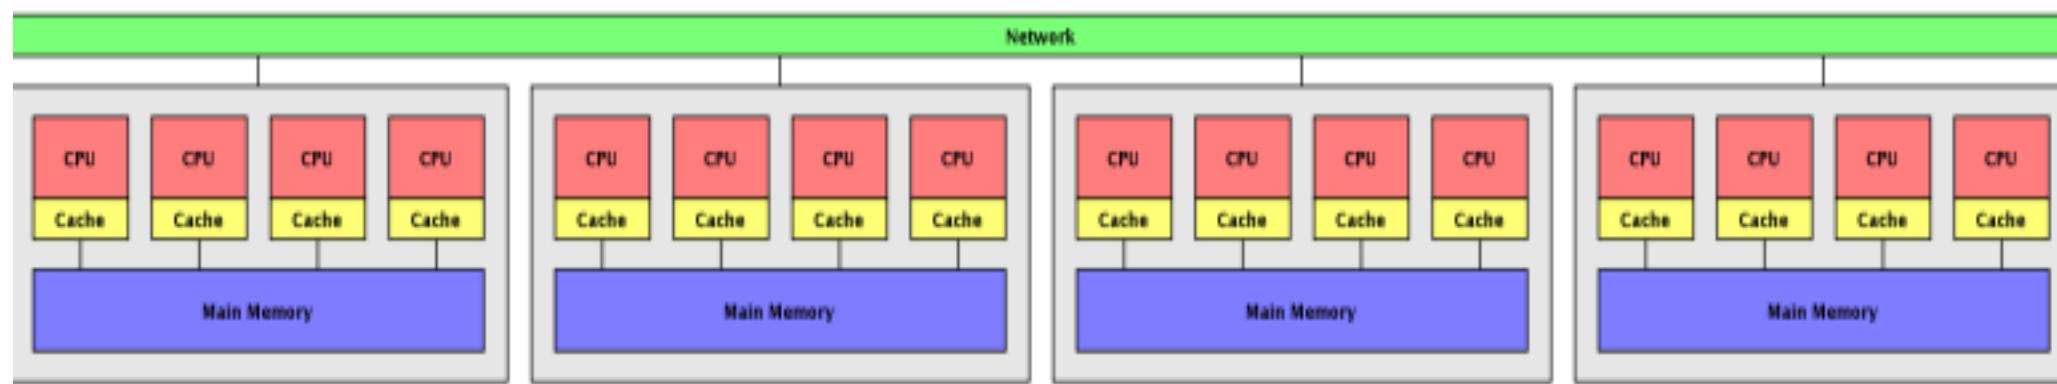
Shared Memory Model

# Simplified Parallel Machine Models



Distributed Memory Model

# Simplified Parallel Machine Models



Hybrid Model

# Writing Programs to Run on Parallel Machines

- C Programming Libraries Exist that provides the programmer an API for writing programs that will run in parallel.
- They provide a **Programming Model** that is portable across architectures, i.e. can provide a message passing model that runs on a shared memory machine.
- We will look at 2 of these Programming Models and Libraries that support the model:
  - Message Passing Programming using MPI (message passing interface)
  - Thread Programming using OpenMP
- As will all libraries they can incur an overhead.

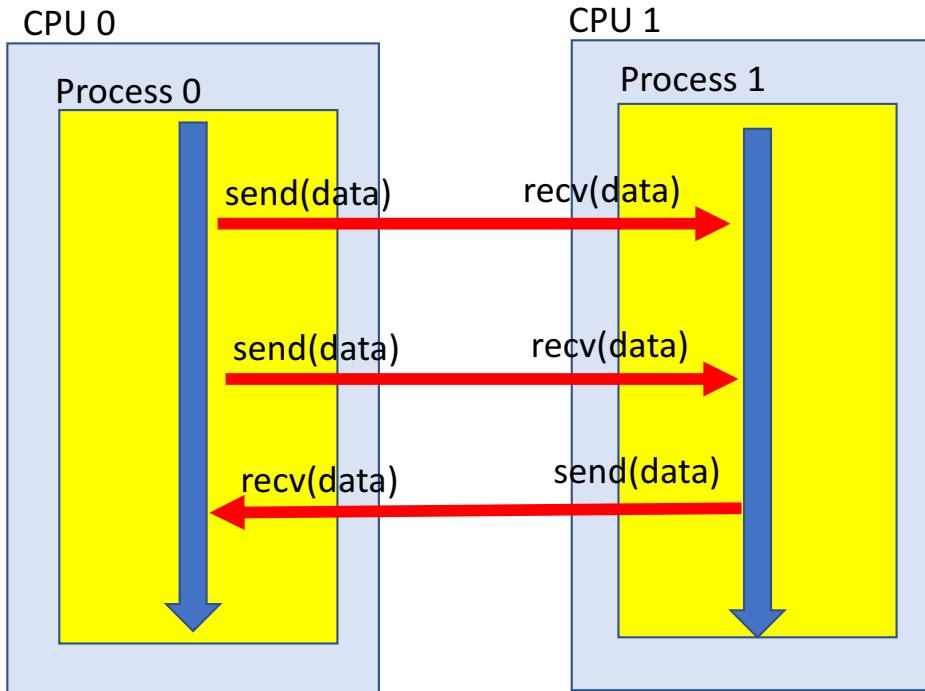
# Outline

- Parallel Machines & Parallel Machine Models
- Parallel Programming
  - Message Passing With MPI
  - Shared Memory Programming with OpenMP

**Ignoring co-processors and GPUs**

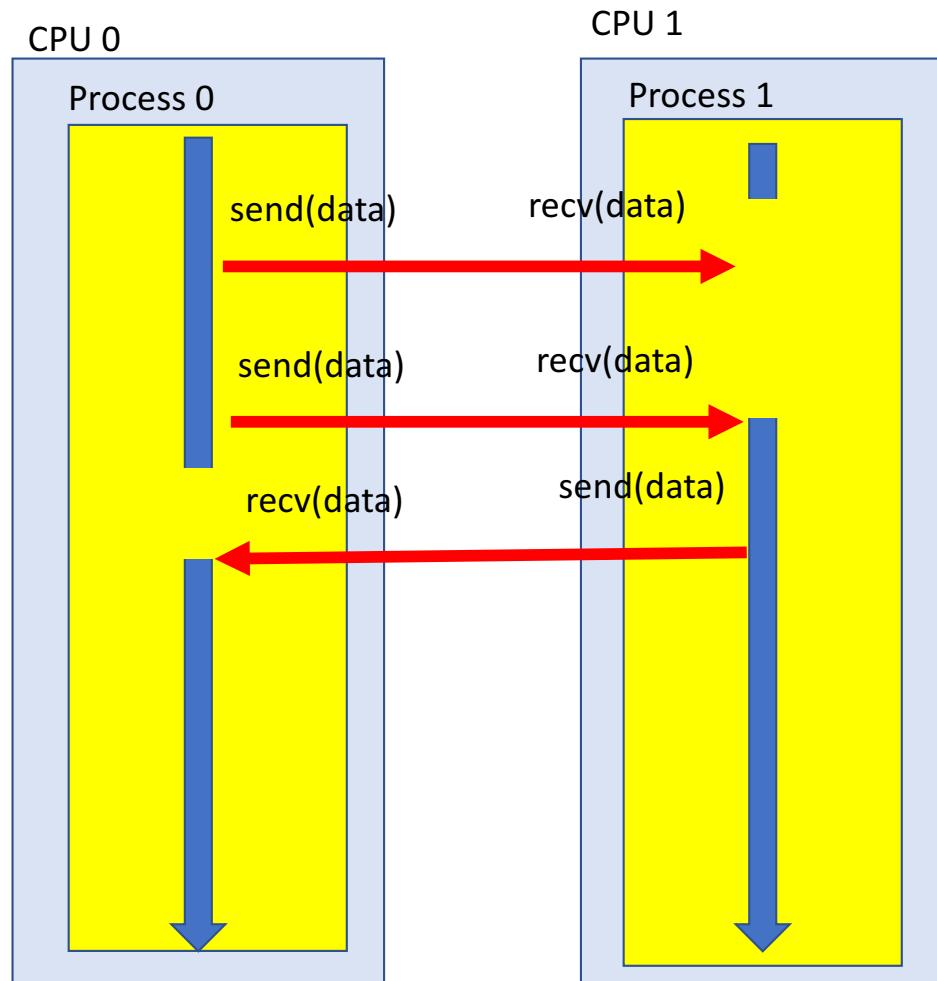
# Message Passing Model

- Processes run independently in their own memory space and processes communicate with each other when data needs to be shared

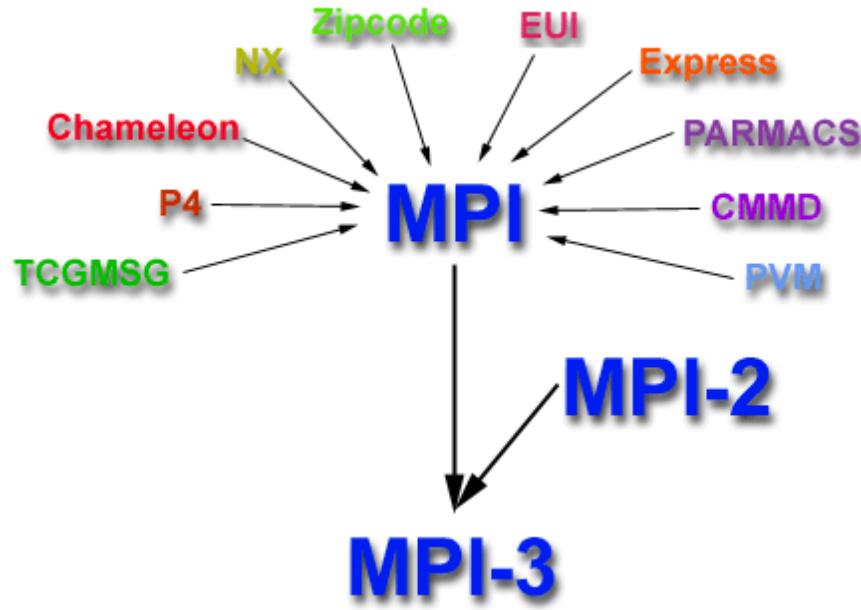


- Basically you write sequential applications with additional function calls to send and recv data.

# Only get Speedup if processes can be kept busy



# Programming Libraries



- Coalasced around a single standard MPI
- Allows for portable code

# MPI

Provides a number of **functions**:

## 1. Enquiries

- How many processes?
- Which one am I?
- Any messages Waiting?

## 2. Communication

- Pair-wise point to point send and receive
- Collective/Group: Broadcast, Scatter/Gather
- Compute and Move: sum, product, max ...

## 3. Synchronization

- Barrier

REMEMBER:

What I am about to show is just C code with some functions you have not yet seen.

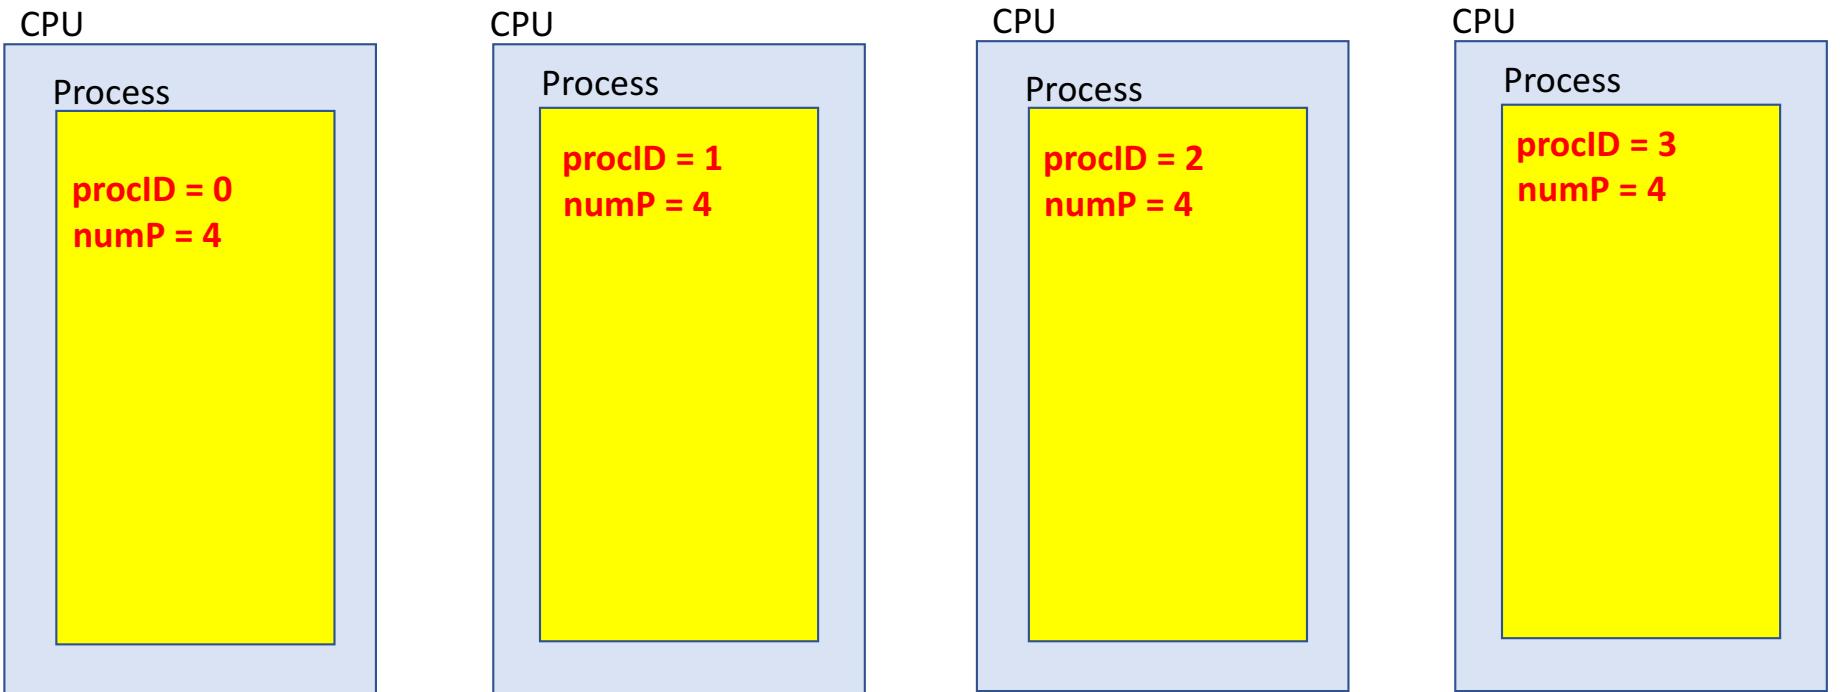
# Hello World

Code/Parallel/mpi/hello1.c

```
#include <mpi.h>
#include <stdio.h>          MPI functions (and MPI_COMM_WORLD) are defined in mpi.h

int main( int argc, char **argv)
{
    int procID, numP;          MPI_Init() and MPI_finalize() must be first and last functions called
                                MPI_COMM_WORLD is a default group containing all processes

    MPI_Init( &argc, &argv );          MPI_Comm_size returns # of
    MPI_Comm_size( MPI_COMM_WORLD, &numP );  processes in the group
    MPI_Comm_rank( MPI_COMM_WORLD, &procID );  MPI_Comm_rank returns processes
                                                unique ID the group, 0 through
printf( "Hello World, I am %d of %d\n", procID, numP );  (numP-1)
MPI_Finalize();
return 0;
}
```



# Send/Recv blocking

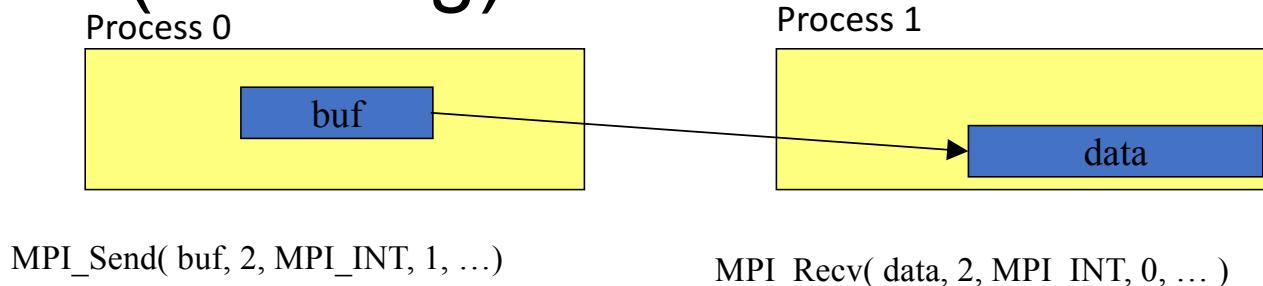
Code/Parallel/mpi/send1.c

```
#include <mpi.h>
#include <stdio.h>
int main( int argc, char **argv) {
    int procID;
    MPI_Status status;
    MPI_Init(&argv, &argc);
    MPI_Comm_rank( MPI_COMM_WORLD, &procID );

    if (procID == 0) { // process 0 sends
        int buf[2] = {12345, 67890};
        MPI_Send( &buf, 2, MPI_INT, 1, 0, MPI_COMM_WORLD);
    } NOTE the PAIR of  
Send/Recv
    else if (procID == 1) { // process 1 receives
        int data[2];
        MPI_Recv( &data, 2, MPI_INT, 0, 0, MPI_COMM_WORLD, &status );
        printf( "Received %d %d\n", data[0], data[1] );
    }

    MPI_Finalize();
    return 0;
}
```

# MPI Basic (Blocking) Send

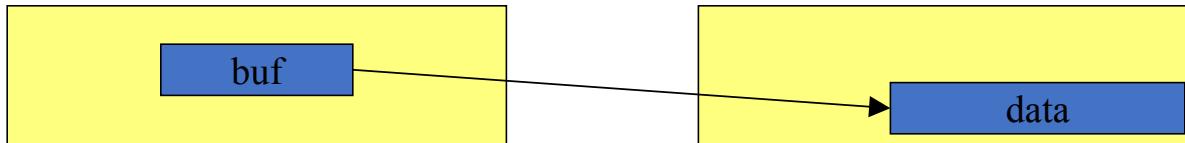


**`MPI_SEND` (*start*, *count*, *datatype*, *dest*, *tag*, *comm*)**

- The message buffer is described by **(*start*, *count*, *datatype*)**.
- The target process is specified by ***dest***, which is the rank of the target process in the communicator specified by ***comm***. The message has an identifier ***tag***.
- When this function returns, the data has been delivered to the system and the buffer can be reused. The message may not have been received by the target process.

<b>MPI_CHAR</b>	char
<b>MPI_WCHAR</b>	wchar_t - wide character
<b>MPI_SHORT</b>	signed short int
<b>MPI_INT</b>	signed int
<b>MPI_LONG</b>	signed long int
<b>MPI_LONG_LONG_INT</b> <b>MPI_LONG_LONG</b>	signed long long int
<b>MPI_SIGNED_CHAR</b>	signed char
<b>MPI_UNSIGNED_CHAR</b>	unsigned char
<b>MPI_UNSIGNED_SHORT</b>	unsigned short int
<b>MPI_UNSIGNED</b>	unsigned int
<b>MPI_UNSIGNED_LONG</b>	unsigned long int
<b>MPI_UNSIGNED_LONG_LONG</b>	unsigned long long int
<b>MPI_FLOAT</b>	float
<b>MPI_DOUBLE</b>	double
<b>MPI_LONG_DOUBLE</b>	long double
<b>MPI_C_COMPLEX</b> <b>MPI_C_FLOAT_COMPLEX</b>	float _Complex
<b>MPI_C_DOUBLE_COMPLEX</b>	double _Complex
<b>MPI_C_LONG_DOUBLE_COMPLEX</b>	long double _Complex
<b>MPI_C_BOOL</b>	_Bool
<b>MPI_INT8_T</b> <b>MPI_INT16_T</b> <b>MPI_INT32_T</b> <b>MPI_INT64_T</b>	int8_t int16_t int32_t int64_t
<b>MPI_UINT8_T</b> <b>MPI_UINT16_T</b> <b>MPI_UINT32_T</b> <b>MPI_UINT64_T</b>	uint8_t uint16_t uint32_t uint64_t
<b>MPI_BYTE</b>	8 binary digits
<b>MPI_PACKED</b>	data packed or unpacked with MPI_Pack() / MPI_Unpack

# MPI Basic (Blocking) Receive



`MPI_Send( buf, 2, MPI_INT, 1, ... )`

`MPI_Recv( data, 2, MPI_INT, 0, ... )`

**`MPI_RECV(start, count, datatype, source, tag, comm, status)`**

- Waits until a matching (both **source** and **tag**) message is received from the system, and the buffer can be used
- **source** is rank in communicator specified by **comm**, or **MPI\_ANY\_SOURCE**
- **tag** is a tag to be matched or **MPI\_ANY\_TAG**
- receiving fewer than **count** occurrences of **datatype** is OK, but receiving more is an error
- **status** contains further information (e.g. size of message)

# Retrieving Further Information

- **Status** is a data structure allocated in the user's program.
- In C:

```
int recvd_tag, recvd_from, recvd_count;
MPI_Status status;
MPI_Recv(..., MPI_ANY_SOURCE, MPI_ANY_TAG, ..., &status )
recvd_tag = status.MPI_TAG;
recvd_from = status.MPI_SOURCE;
MPI_Get_count( &status, datatype, &recvd_count );
```

# Not Quite So Simple as ensuring PAIRS of send/recv

```
#include <mpi.h>
#include <stdio.h>
#define DATA_SIZE 1000
int main(int argc, char **argv) {
    int procID, numP;
    MPI_Status status;
    int buf[DATA_SIZE];

    MPI_Init( &argc, &argv );
    MPI_Comm_size( MPI_COMM_WORLD, &numP );
    MPI_Comm_rank( MPI_COMM_WORLD, &procID );
    if (procID == 0) {
        for (int i=0; i<DATA_SIZE; i++) buf[i]=1+i;
        MPI_Send(&buf, DATA_SIZE, MPI_INT, 1, 0, MPI_COMM_WORLD);
        MPI_Recv(&buf, DATA_SIZE, MPI_INT, 1, 0, MPI_COMM_WORLD, &status);
        printf("%d Received %d %d\n", procID, buf[0], buf[DATA_SIZE-1]);
    } else if (procID == 1) {
        for (int i=0; i<DATA_SIZE; i++) buf[i]=DATA_SIZE-i;
        MPI_Send(&buf, DATA_SIZE, MPI_INT, 0, 0, MPI_COMM_WORLD);
        MPI_Recv(&buf, DATA_SIZE, MPI_INT, 0, 0, MPI_COMM_WORLD, &status);
        printf("%d Received %d %d\n", procID, buf[0], buf[DATA_SIZE-1]);
    }
    MPI_Finalize();
    return 0;
}
```

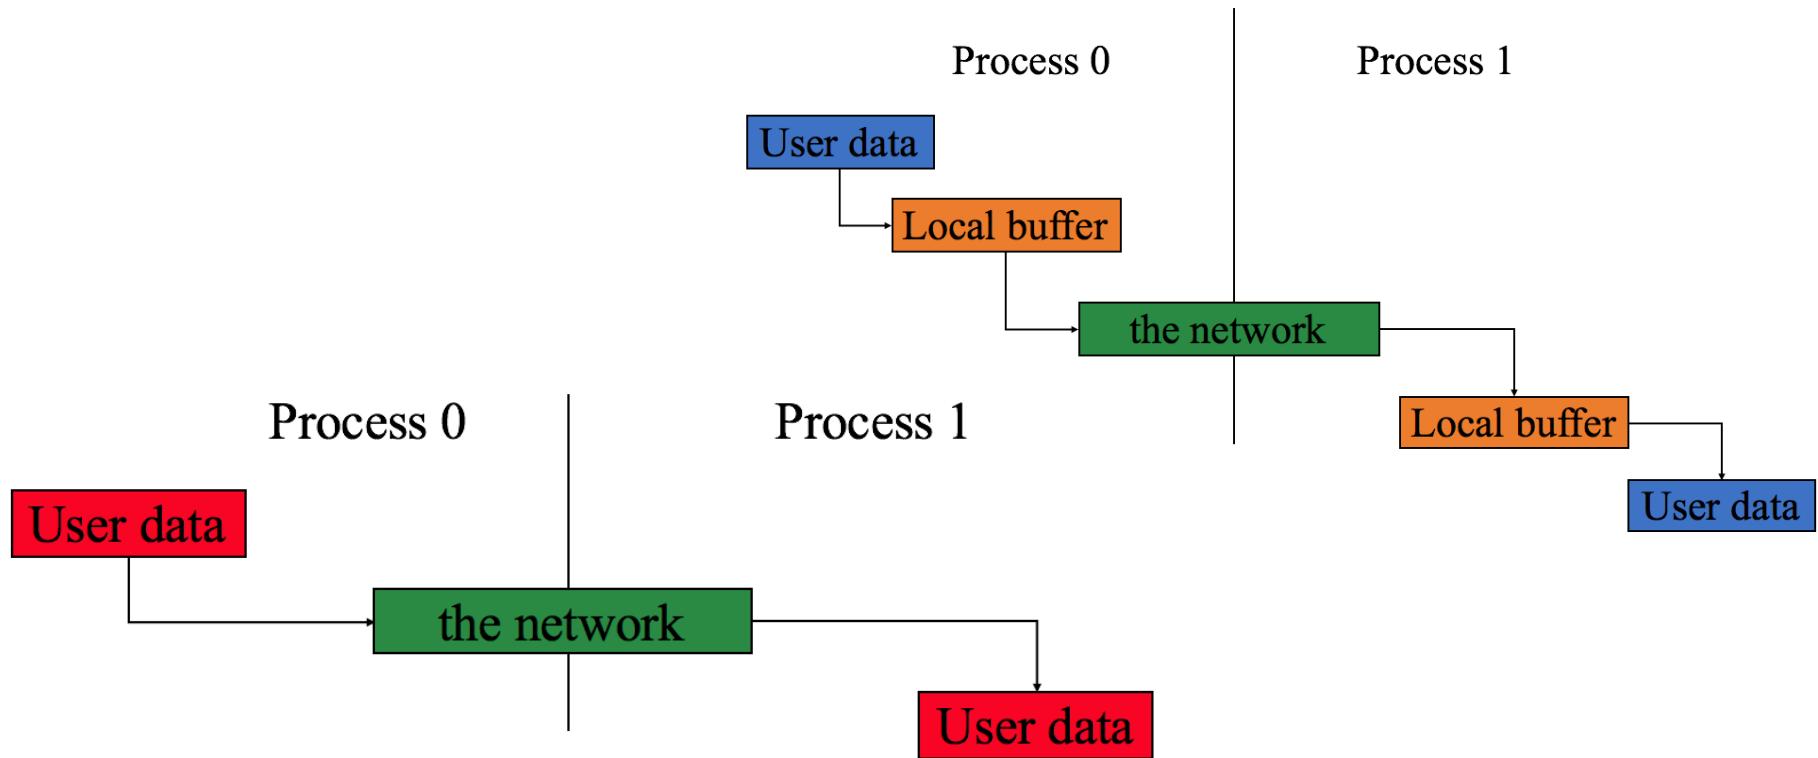
Code/Parallel/mpi/send2.c

```
[mpi >mpicc send2.c; mpirun -n 2 ./a.out
Buffer Size: 1000
0 Received 1000 1
1 Received 1 1000]
```

```
[mpi >mpicc send2.c; mpirun -n 2 ./a.out
Buffer Size: 10000
^Cmpi >
mpi >]
```

**DEADLOCK .. PROGRAM HANGS .. WHY?**

# Why Deadlock? .. Where Does the Data Go



If large message & insufficient data, the send() must wait for buffer to clear through a recv()

source: CS267, Jim Demmel

## Current Problem:

Process 0	Process 1
<b>Send(1)</b>	<b>Send(0)</b>
<b>Recv(1)</b>	<b>Recv(0)</b>

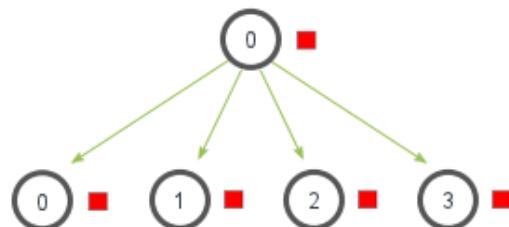
Could revise order  
this requires some code rewrite:

Process 0	Process 1
<b>Send(1)</b>	<b>Recv(0)</b>
<b>Recv(1)</b>	<b>Send(0)</b>

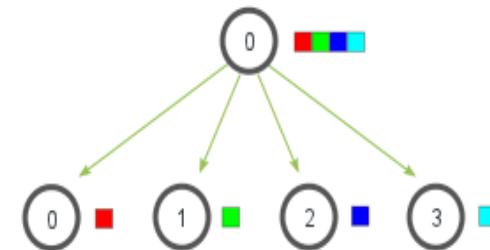
Alternatives use non-blocking sends.

# Some Collective Functions

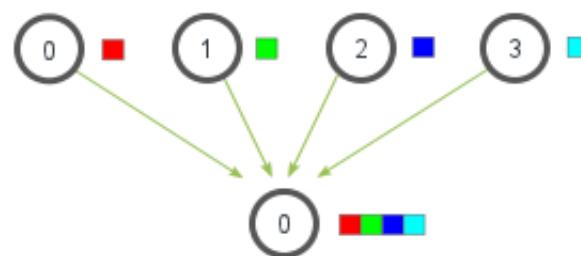
MPI\_Bcast



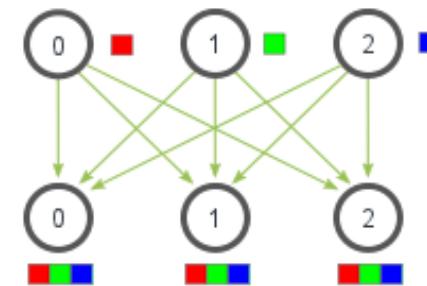
MPI\_Scatter



MPI\_Gather



MPI\_Allgather



# Broadcast

```
#include <mpi.h>
#include <stdio.h>

int main( int argc, char **argv) {
    int procID; buf[2];

    MPI_Init( &argc, &argv );
    MPI_Comm_rank( MPI_COMM_WORLD, &procID );

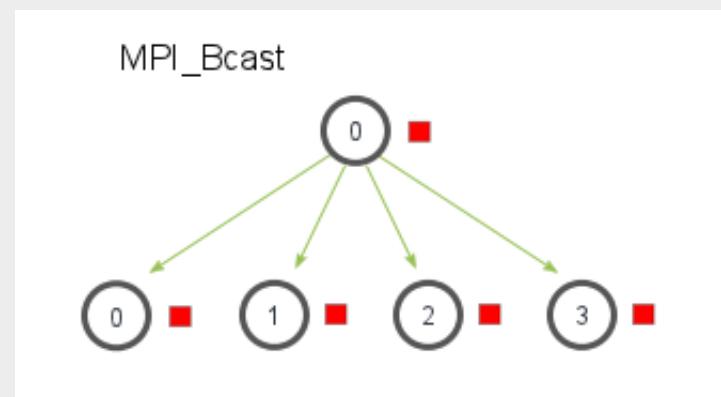
    if (procID == 0) {
        buf[0] = 12345;
        buf[1] = 67890;
    }

    MPI_Bcast(&buf, 2, MPI_INT, 0, MPI_COMM_WORLD);

    printf("Process %d data %d %d\n", procID, buf[0], buf[1]);

    MPI_Finalize();
    return 0;
}
```

Code/Parallel/mpi/bcast1.c



# Scatter

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>

#define LUMP 5

int main(int argc, char **argv) {
    int numP, procID;
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &numP);
    MPI_Comm_rank(MPI_COMM_WORLD, &procID);

    int *globalData=NULL;
    int localData[LUMP];

    if (procID == 0) { // malloc and fill in with data
        globalData = malloc(LUMP * numP * sizeof(int) );
        for (int i=0; i<LUMP*numP; i++)
            globalData[i] = i;
    }

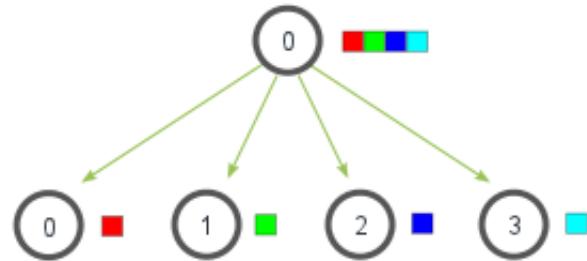
    MPI_Scatter(globalData, LUMP, MPI_INT, &localData, LUMP, MPI_INT, 0, MPI_COMM_WORLD);
    printf("Processor %d has first: %d last %d\n", procID, localData[0], localData[LUMP-1]);

    if (procID == 0) free(globalData);

    MPI_Finalize();
}
```

Code/Parallel/mpi/scatter1.c

MPI\_Scatter



```

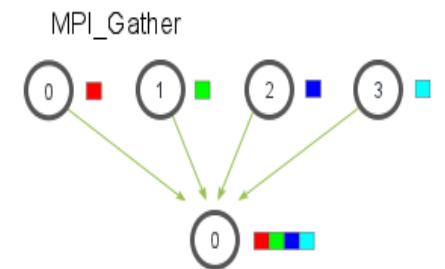
#include "mpi.h"
#include <stdio.h>
#define LUMP 5
int main(int argc, const char **argv) {
    int procID, numP, ierr;
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &numP);
    MPI_Comm_rank(MPI_COMM_WORLD, &procID);

    int *globalData=NULL;
    int localData[LUMP];
    if (procID == 0) { // malloc global data array only on P0
        globalData = malloc(LUMP * numP * sizeof(int));
    }
    for (int i=0; i<LUMP; i++)
        localData[i] = procID*10+i;

    MPI_Gather(localData, LUMP, MPI_INT, globalData, LUMP, MPI_INT, 0, MPI_COMM_WORLD);

    if (procID == 0) {
        for (int i=0; i<numP*LUMP; i++)
            printf("%d ", globalData[i]);
        printf("\n");
    }
    if (procID == 0) free(globalData);
    MPI_Finalize();
}

```

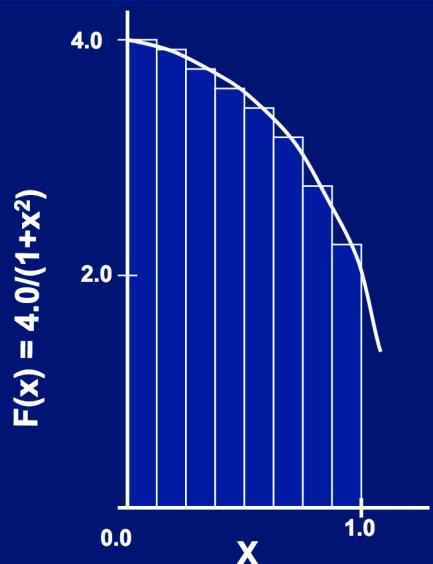


# MPI can be simple

- Claim: most MPI applications can be written with only 6 functions (although which 6 may differ)
- Using point-to-point:
  - `MPI_INIT`
  - `MPI_FINALIZE`
  - `MPI_COMM_SIZE`
  - `MPI_COMM_RANK`
  - `MPI_SEND`
  - `MPI_RECEIVE`
- Using collectives:
  - `MPI_INIT`
  - `MPI_FINALIZE`
  - `MPI_COMM_SIZE`
  - `MPI_COMM_RANK`
  - `MPI_BCAST/MPI_SCATTER`
  - `MPI_GATHER/MPI_ALLGATHER`
- You may use more for convenience or performance

# Exercise: Parallelize Compute PI using MPI

## Numerical Integration



Mathematically, we know that:

$$\int_0^1 \frac{4.0}{(1+x^2)} dx = \pi$$

We can approximate the integral as a sum of rectangles:

$$\sum_{i=0}^N F(x_i) \Delta x \approx \pi$$

Where each rectangle has width  $\Delta x$  and height  $F(x_i)$  at the middle of interval  $i$ .

```
#include <stdio>
static int long numSteps = 100000;
int main() {
    double pi = 0; double time=0;
    // your code
    for (int i=0; i<numSteps; i++) {
        // your code
    }
    // your code
    printf("PI = %f, duration: %f ms\n",pi, time);
    return 0;
}
```

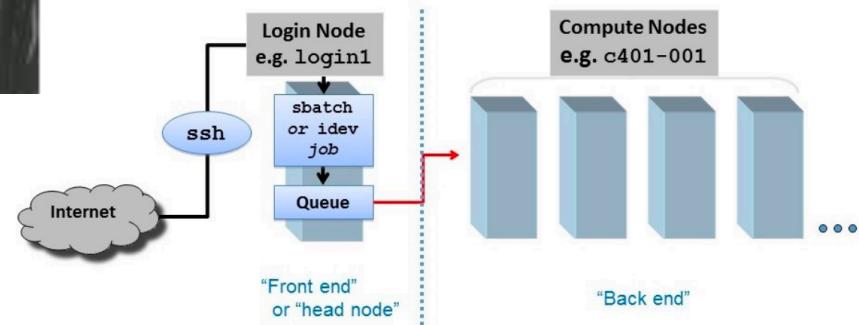
Source: UC Berkeley, Tim Mattson (Intel Corp), CS267 & elsewhere

# Stampede2



TACC

DESIGNSAFE-CI



# sbatch script

```
#!/bin/bash
#-----
# Generic SLURM script – MPI Hello World
#
# This script requests 1 node and 8 cores/node (out of total 64 avail)
# for a total of 1*8 = 8 MPI tasks.
#-----
#SBATCH -J myjob      # Job name
#SBATCH -o myjob.%j.out # stdout; %j expands to jobid
#SBATCH -e myjob.%j.err # stderr; skip to combine stdout and stderr
#SBATCH -p development  # queue
#SBATCH -N 1          # Number of nodes, not cores (64 cores/node)
#SBATCH -n 8          # Total number of MPI tasks (if omitted, n=N)
#SBATCH -t 00:00:10    # max time

module petsc  # load any needed modules, these just examples
module load list

ibrun ./a.out
```

1. Put your pi.c into your github repository and push it to your fork

```
cp pi.c ~/SimCenterBootcamp/Code/Parallel/mpi/pi.c  
cd ~/SimCenterBootcamp/Code/Parallel/mpi  
git add pi.c  
git commit -m "pi.c initial import"  
git push
```

2. Login to stampede2

```
ssh yourlogin@stampede2.tacc.utexas.edu
```

3. Now on stampede2 login node, clone the repository on stampede2

```
git clone http://???????
```

4. cd to the ~/SimCenterBootCamp/Code/Parallel/mpi directory

```
cd ~/SimCenterBootcamp/Code/Parallel/mpi
```

5. Compile hello1.c or whatever you want

```
mpicc hello1.c
```

6. Submit the job to SLURM queue

```
sbatch submit.sh
```

7. Look at the output

```
cat myjob.*.out
```

8. Make changes to your code, compile & run!

```
Your on your own
```

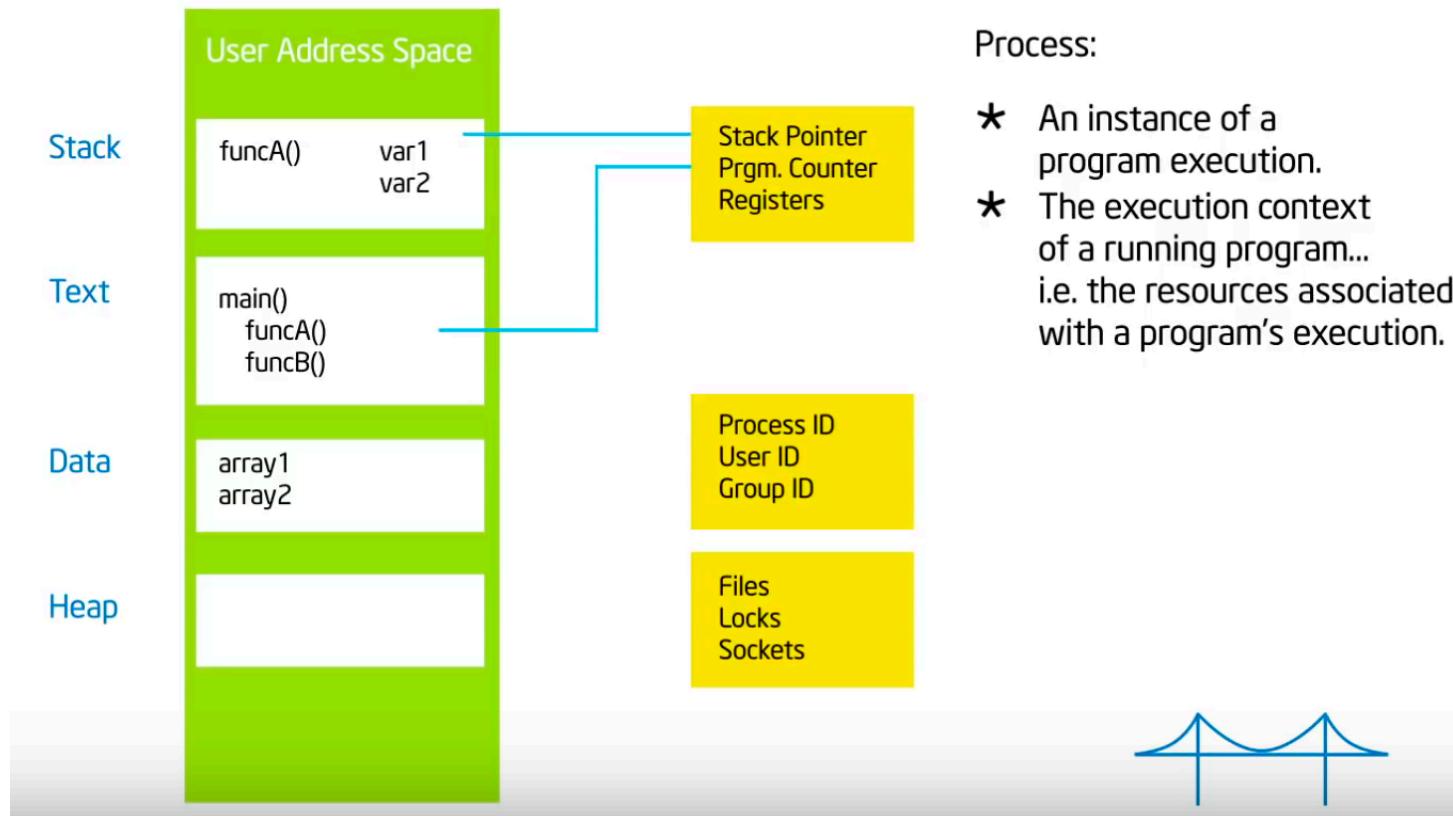
# SLURM script

# Outline

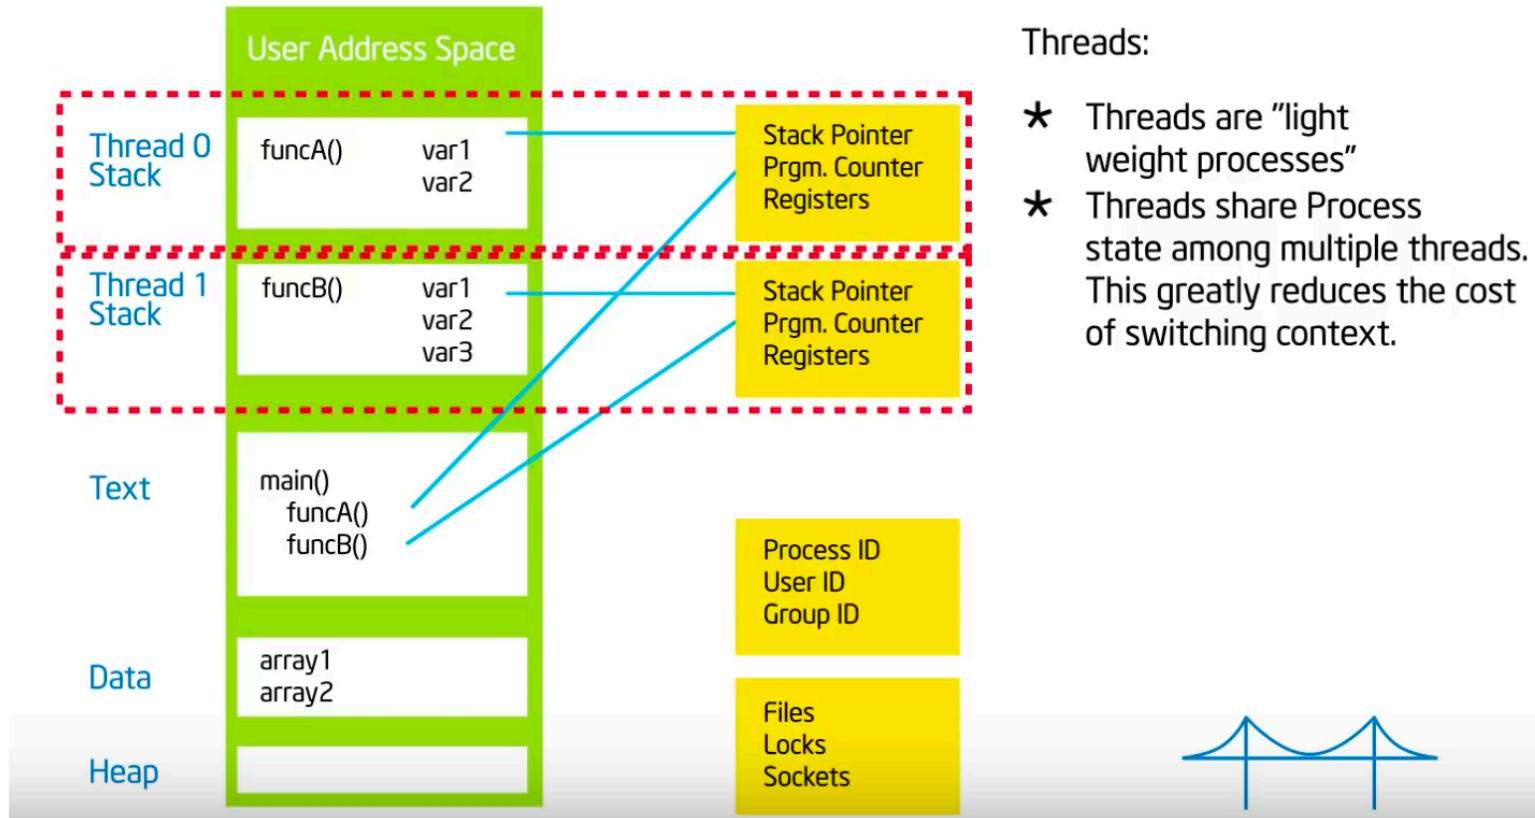
- Parallel Machines & Parallel Machine Models
- Parallel Programming
  - Message Passing With MPI
  - Shared Memory Programming with OpenMP

**Ignoring co-processors and GPUs**

# Process

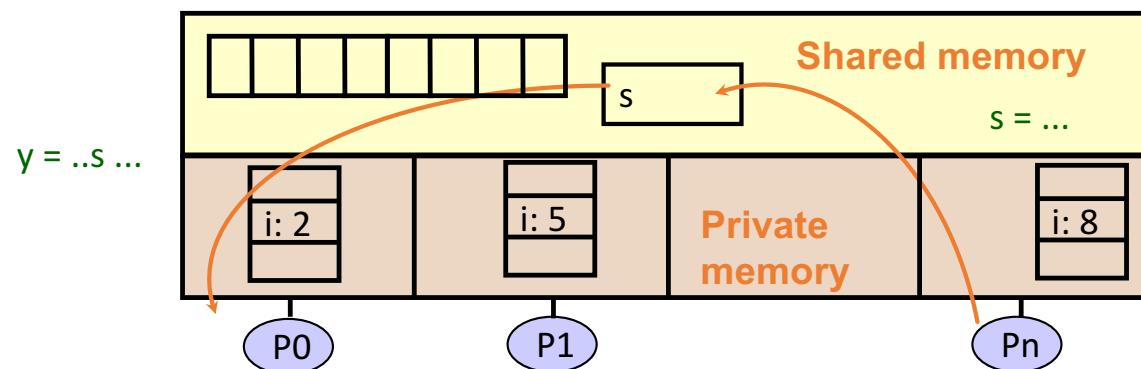


# Threads



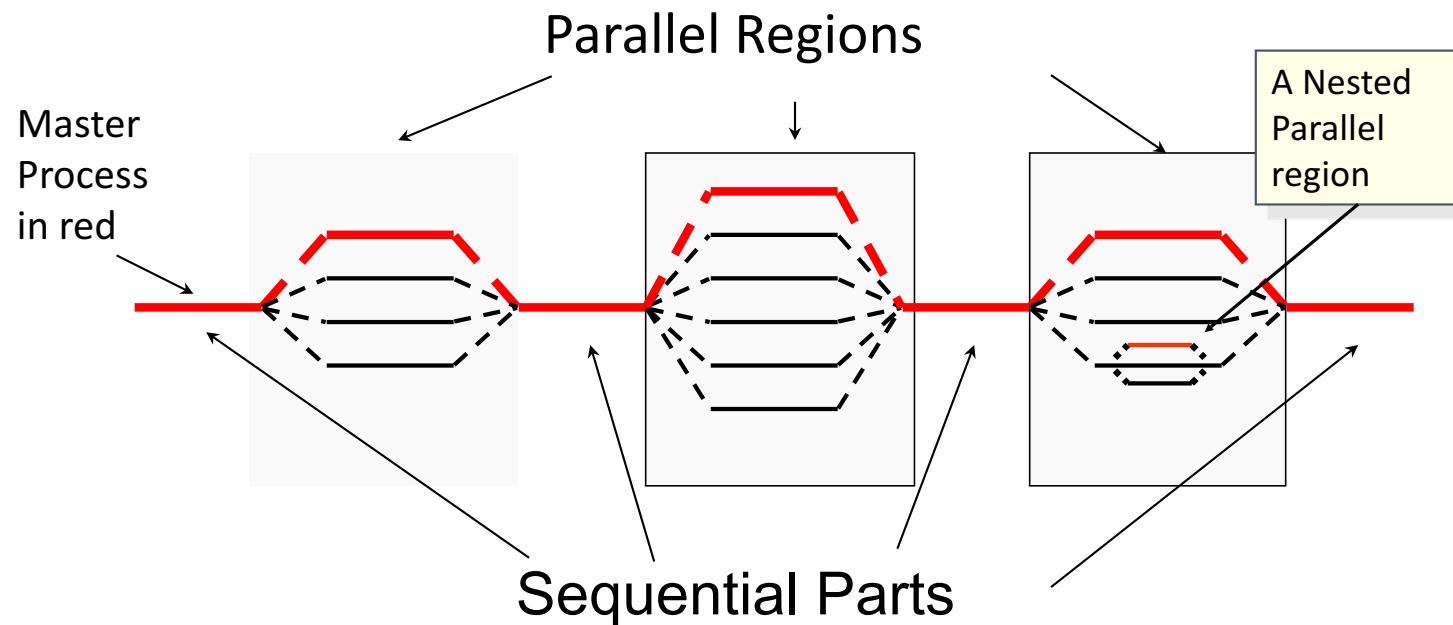
# Threads

- Can be created dynamically, mid-execution, in some languages
- Each thread has a set of **private variables**, e.g., local stack variables
- Also a set of **shared variables**, e.g., static variables, shared common blocks, or global heap.
  - Threads communicate **implicitly** by writing and reading shared variables.
  - Threads coordinate by **synchronizing** on shared variables



## Programming Model for Threads

- Master Process spawns a team of threads as needed.
- Parallelism added incrementally until performance goals are met, i.e., the sequential program evolves into a parallel program.



# Runtime Library Options for Shared Memory

POSIX Threads (pthreads)

OpenMP



- OpenMP provides multi-threaded capabilities to C, C++ and Fortran Programs
- In a threaded environment, threads communicate by sharing data
- Unintended sharing of data causes **race conditions**
- **Race Condition:** program output is different every time you run the program, a consequence of the threads being scheduled differently
- OpenMP provides constructs to control what blocks of code are run in parallel and also constructs for providing access to shared data using synchronization
- Synchronization has overhead consequences, you have to minimize them to get good speedup.

- Mostly Set of Compiler directives (#pragma) applying to structured block

```
#pragma omp parallel
```

- Some runtime library calls

```
omp_num_threads(4);
```

- Being compiler directives, they are built into most compilers .
- Just have to activate it when compiling

```
gcc hello.c -fopenmp
```

```
icc hello.c /Qopenmp
```

# Hello World

```
#include <omp.h>
#include <stdio.h>
```

```
int main( int argc, char *argv[] )
{
    #pragma omp parallel
    {
        int id = omp_get_thread_num();
        int numP = omp_get_num_threads();
        printf("Hello World, I am %d of %d\n",id,numP);
    }
    return 0;
}
```

```
openmp >gcc-7.2 hello1.c -fopenmp; ./a.out
Hello World, I am 0 of 4
Hello World, I am 3 of 4
Hello World, I am 1 of 4
Hello World, I am 2 of 4
openmp >
```

```
openmp >export env OMP_NUM_THREADS=2
openmp >./a.out
Hello World, I am 0 of 2
Hello World, I am 1 of 2
openmp >
```

Code/Parallel/openmp/hello1.c

# Hello World – changing num threads

```
#include <openmp.h>
#include <stdio.h>
```

Code/Parallel/openmp/hello2.c

```
int main( int argc, char *argv[] )
{
    omp_set_num_threads(2);
    #pragma omp parallel
    {
        int id = omp_get_thread_num();
        int numP = omp_get_num_threads();
        printf("Hello World, I am %d of %d\n");
    }
    return 0;
}
```

Runtime function to  
request a certain  
number of threads

Runtime function to  
return actual number  
of threads in the  
team

```
#include <openmp.h>
#include <stdio.h>

int main( int argc, char *argv[] )
{
#pragma omp parallel num_threads(2)
{
    int id = omp_get_thread_num();
    int numP = omp_get_num_threads();
    printf("Hello World, I am %d of %d\n",id,numP);
}
return 0;
}
```

Code/Parallel/openmp/hello3.c

# Different # threads in different blocks

```
#include <omp.h>
#include <stdio.h>

int main(int argc, const char **argv) {

#pragma omp parallel num_threads(2)
{
    int id = omp_get_thread_num();
    int numP = omp_get_num_threads();
    printf("Hello World, I am %d of %d\n",id,numP);
}

#pragma omp parallel num_threads(4)
{
    int id = omp_get_thread_num();
    int numP = omp_get_num_threads();
    printf("Hello World Again, I am %d of %d\n",id,numP);
}

return(0);
}
```

Code/Parallel/openmp/hello4.c

```
openmp >gcc-7.2 hello4.c -fopenmp; ./a.out
Hello World, I am 0 of 2
Hello World, I am 1 of 2
Hello World Again, I am 1 of 4
Hello World Again, I am 2 of 4
Hello World Again, I am 3 of 4
Hello World Again, I am 0 of 4
openmp >
```

# Hello World – shared variables & RACE CONDITIONS

```
#include <omp.h>
#include <stdio.h>

int main(int argc, const char **argv) {
    int id, numP;

#pragma omp parallel num_threads(4)
{
    id = omp_get_thread_num();
    numP = omp_get_num_threads();
    printf("Hello World from %d of %d threads\n", id, numP);
}

return(0);
}
```

Code/Parallel/openmp/hello5.c

```
openmp >gcc-7.2 hello5.c -fopenmp; ./a.out
Hello World from 1 of 4 threads
Hello World from 2 of 4 threads
Hello World from 3 of 4 threads
Hello World from 0 of 4 threads
openmp >gcc-7.2 hello5.c -fopenmp; ./a.out
Hello World from 0 of 4 threads
Hello World from 3 of 4 threads
Hello World from 2 of 4 threads
Hello World from 1 of 4 threads
openmp >gcc-7.2 hello5.c -fopenmp; ./a.out
Hello World from 0 of 4 threads
Hello World from 0 of 4 threads
Hello World from 2 of 4 threads
Hello World from 3 of 4 threads
```

# What We Want Threads To DO

- Work Independently With Controlled Access at times to Shared Data
  - Parallel Tasks Constructs
    - `omp parallel`
    - `Opf for`
  - Shared Data

# Simple Vector Sum

```
#include <omp.h>
#include <stdio.h>
#define DATA_SIZE 10000

void sumVectors(int N, double *A, double *B, double *C, int tid, int numT);
|  
nt main(int argc, const char **argv) {
    double a[DATA_SIZE], b[DATA_SIZE], c[DATA_SIZE];
    int num;
    for (int i=0; i<DATA_SIZE; i++) { a[i] = i+1; b[i] = i+1; }
    double tdata = omp_get_wtime();
#pragma omp parallel
{
    int tid = omp_get_thread_num();
    int numT = omp_get_num_threads();
    num = numT;
    sumVectors(DATA_SIZE, a, b, c, tid, numT);
}
    tdata = omp_get_wtime() - tdata;
    printf("first %f last %f in time %f using %d threads\n"
    return 0;
}
```

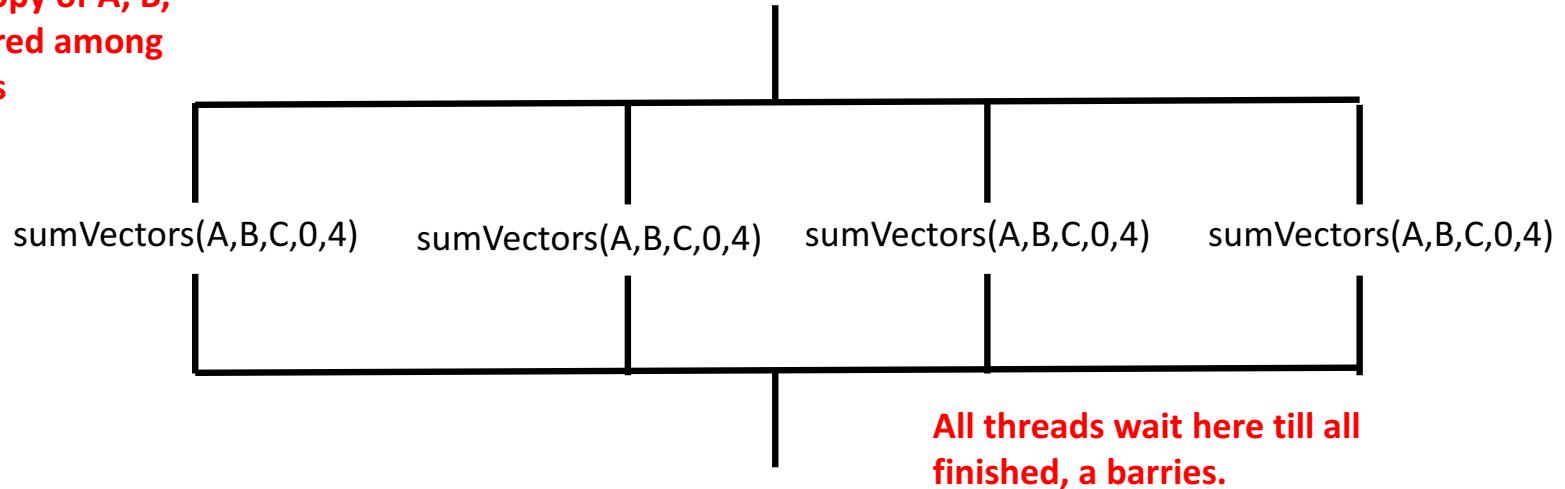
Code/Parallel/openmp/sum1.c

```
void sumVectors(int N, double *A, double *B, double *C, int tid, int numT)
    // determine start & end for each thread
    int start =  tid * N / numT;
    int end = (tid+1) * N / numT;
    if (tid == numT-1)
        end = N;

    // do the vector sum for threads bounds
    for(int i=start; i<end; i++) {
        C[i] = A[i]+B[i];
    }
}
```

# Implicit Barrier in Code

A single copy of A, B,  
and C shared among  
all threads



```
openmp >export env OMP_NUM_THREADS=1; ./a.out
first 2.000000 last 200000.000000 in time 0.000902 using 1 threads
openmp >export env OMP_NUM_THREADS=2; ./a.out
first 2.000000 last 200000.000000 in time 0.000678 using 2 threads
openmp >export env OMP_NUM_THREADS=4; ./a.out
first 2.000000 last 200000.000000 in time 0.000652 using 4 threads
openmp >export env OMP_NUM_THREADS=8; ./a.out
first 2.000000 last 200000.000000 in time 0.000693 using 8 threads
```

# The for is such an obvious candidate for threads:

```
#include <omp.h>
#include <stdio.h>
#include <math.h>
#define DATA_SIZE 10000

int main(int argc, const char **argv) {
    double a[DATA_SIZE], b[DATA_SIZE], c[DATA_SIZE];
    for (int i=0; i<DATA_SIZE; i++) { a[i] = i+1; b[i] = i+1; }
    double tdata = omp_get_wtime();
#pragma omp parallel
{
    #pragma omp for
    for (int i=0; i<DATA_SIZE; i++)
        c[i] = a[i]+b[i];
}
tdata = omp_get_wtime() - tdata;
printf("first %f last %f in time %f \n",c[0], c[DATA_SIZE-1], tdata);
return 0;
}
```

Code/Parallel/openmp/sum2.c

Code/Parallel/openmp/sum3.c

```
#pragma omp parallel for
for (int i=0; i<DATA_SIZE; i++)
    c[i] = a[i]+b[i];
```

# How About Dot Product?

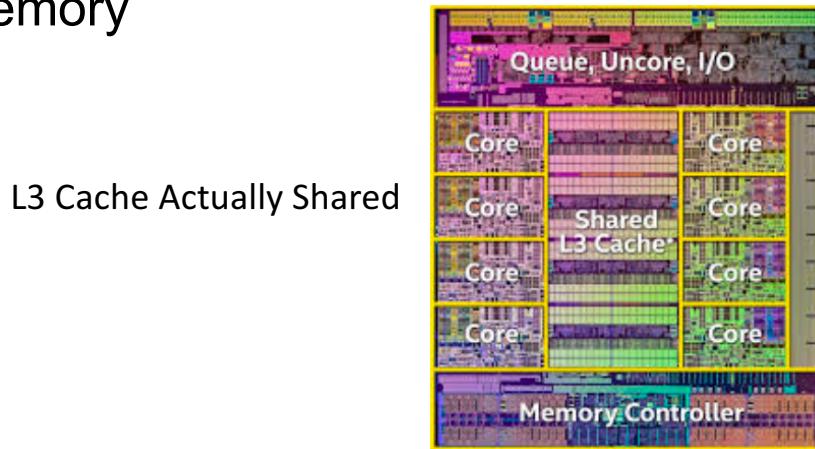
```
#include <omp.h>           Code/Parallel/openmp/dot1.c
#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#define DATA_SIZE 10000

int main(int argc, const char **argv) {
    int nThreads = 0;
    double dot = 0, a[DATA_SIZE], sum[64];          Create a shared array to store data
    for (int i=0; i<DATA_SIZE; i++) a[i] = i+1;
    for (int i=0; i<64; i++) sum[i] = 0;

#pragma omp parallel
{
    int tid = omp_get_thread_num();
    int numT = omp_get_num_threads();
    if (tid == 0) nThreads = numT;
    for (int i=tid; i<DATA_SIZE; i+= numT)
        sum[tid] += a[i]*a[i];
}
for (int i=0; i<nThreads; i++)
    dot += sum[i];          Combine sequentially
dot = sqrt(dot);
printf("dot %f\n", dot);
return 0;
}
```

# Poor Performance?

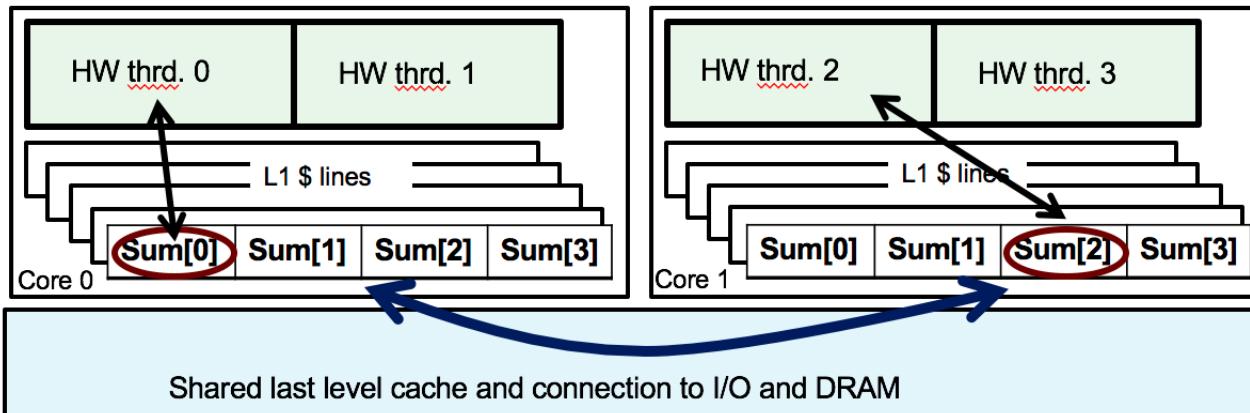
- Want high performance for shared memory: Use Caches!
  - Each processor has its own cache (or multiple caches)
  - Place data from memory into cache
  - Writeback cache: don't send all writes over bus to memory



- Problem is in multi-threaded model with all threads wanting to WRITE same spatially temporal data we have contention at the cache line in the L3 cache

# False Sharing

- If independent data elements happen to sit on the same cache line, each update will cause the cache lines to “slosh back and forth” between threads ...  
This is called **false sharing** or sequential **consistency**.



- Sequential Consistency problem is pervasive and performance critical in shared memory

Solution?

source: UC Berkeley, Tim Mattson (Intel Corp), CS267 & elsewhere

# AVOID FALSE SHARING

```
#include <omp.h>
#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#define DATA_SIZE 10000
#define PAD 64

int main(int argc, const char **argv) {
    int nThreads = 0;
    double dot = 0, a[DATA_SIZE], sum[64][PAD];
    for (int i=0; i<DATA_SIZE; i++) a[i] = i+1;
    for (int i=0; i<64; i++) sum[i][0]= 0;

#pragma omp parallel
{
    int tid = omp_get_thread_num();
    int numT = omp_get_num_threads();
    if (tid == 0) nThreads = numT;
    for (int i=tid; i<DATA_SIZE; i+= numT)
        sum[tid][0]+= a[i]*a[i];
}
    for (int i=0; i<nThreads; i++)
        dot += sum[i][0];
    dot = sqrt(dot);
    printf("dot %f \n", dot);
    return 0;
```

Code/Parallel/openmp/dot2.c

Pad the shared array to store data to avoid false sharing

# SYNCHRONIZATION

```
#include <omp.h>          Code/Parallel/openmp/dot3.c
#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#define DATA_SIZE 10000

int main(int argc, const char **argv) {
    double dot = 0;
    double a[DATA_SIZE];
    for (int i=0; i<DATA_SIZE; i++) a[i] = i+1;

#pragma omp parallel
{
    int tid = omp_get_thread_num();
    int numT = omp_get_num_threads();
    double sum = 0.;
    for (int i=tid; i<DATA_SIZE; i+= numT)
        sum += a[i]*a[i];
#pragma omp critical
    dot += sum;
}
dot = sqrt(dot);
printf("dot %f \n", dot);
return 0;
}
```

# REDUCTION

```
#include <omp.h>          Code/Parallel/openmp/dot4.c
#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#define DATA_SIZE 10000

int main(int argc, const char **argv) {
    double dot = 0;
    double a[DATA_SIZE];
    for (int i=0; i<DATA_SIZE; i++) a[i] = i+1;

#pragma omp parallel reduction(+:dot)
{
    int tid = omp_get_thread_num();
    int numT = omp_get_num_threads();
    double sum = 0.;
    for (int i=tid; i<DATA_SIZE; i+= numT)
        sum += a[i]*a[i];
// #pragma omp critical
    dot += sum;
}
dot = sqrt(dot);
printf("dot %f \n", dot);
return 0;
}
```

# Additional Reduction Operators

Operator	Initial value
+	0
*	1
-	0
min	Largest pos. number
max	Most neg. number

# Are Pitfalls to Parallel Loops

- Basic approach might only occur for experienced programmer!
  - Find compute intensive loops
  - Make the loop iterations independent ... So they can safely execute in any order without loop-carried dependencies
  - Place the appropriate OpenMP directive and test

```
int i, j, A[MAX];
j = 5;
for (i=0;i< MAX; i++) {
    j +=2;
    A[i] = big(j);
}
```

Note: loop index  
“i” is private by  
default

Remove loop  
carried  
dependence

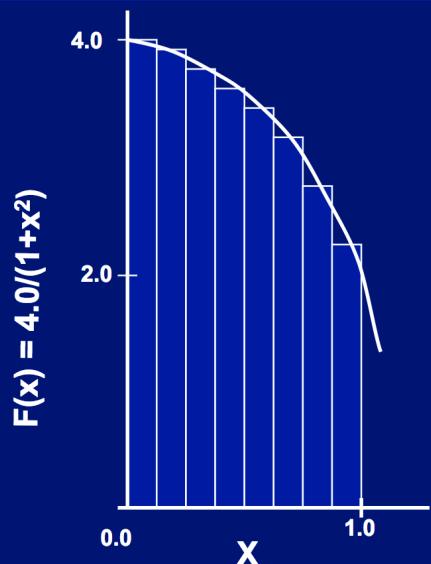
```
int i, A[MAX];
#pragma omp parallel for
for (i=0;i< MAX; i++) {
    int j = 5 + 2*(i+1);
    A[i] = big(j);
}
```



- OpenMP provides multi-threaded capabilities to C, C++ and Fortran Programs
- In a threaded environment, threads communicate by sharing data
- Unintended sharing of data causes **race conditions**
- **Race Condition:** program output is different every time you run the program, a consequence of the threads being scheduled differently
- OpenMP provides constructs to control what blocks of code are run in parallel and also constructs for providing access to shared data using synchronization
- Synchronization has overhead consequences, you have to minimize them to get good speedup.

# Exercise: Parallelize Compute PI using OpenMP

## Numerical Integration



Mathematically, we know that:

$$\int_0^1 \frac{4.0}{(1+x^2)} dx = \pi$$

We can approximate the integral as a sum of rectangles:

$$\sum_{i=0}^N F(x_i) \Delta x \approx \pi$$

Where each rectangle has width  $\Delta x$  and height  $F(x_i)$  at the middle of interval  $i$ .

```
#include <stdio>
static int long numSteps = 100000;
int main() {
    double pi = 0; double time=0;
    // your code
    for (int i=0; i<numSteps; i++) {
        // your code
    }
    // your code
    printf("PI = %f, duration: %f ms\n",pi, time);
    return 0;
}
```

Source: UC Berkeley, Tim Mattson (Intel Corp), CS267 & elsewhere